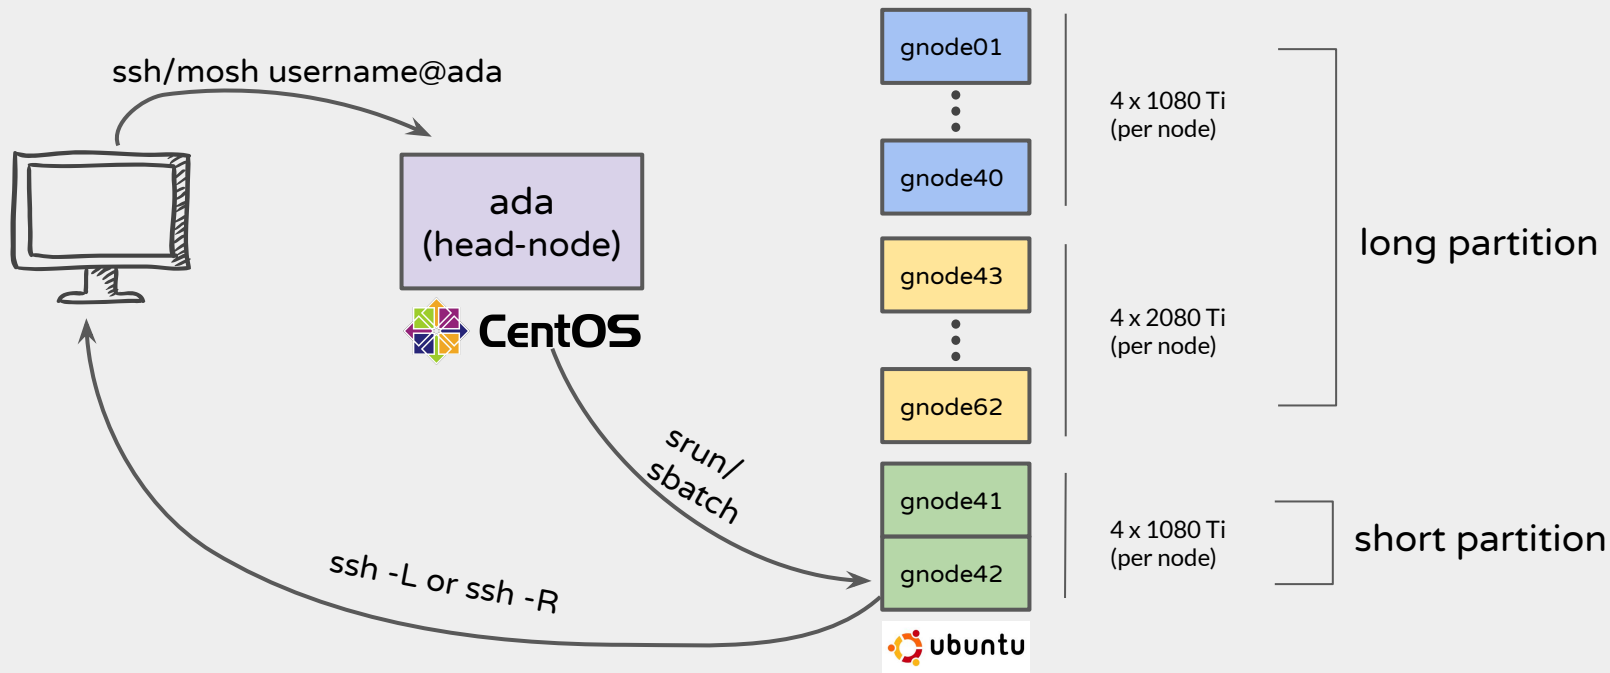


# \$ ada-tutorial

## \$ init\_

- ada is your compute cluster. All your GPU-intensive jobs will run on this.
- tutorial is linux-oriented, entire ada ecosystem is as well.
- cluster is managed by slurm - a job scheduling and cluster management system
- each student has a separate slurm account. Consider it like a bank account with gpu mins as currency. Use it wisely.

# \$ model\_



\$ ~~ssh~~ mosh

> Install

`sudo apt-get install mosh` (Debian)

`brew install mosh` (MacOS)

> Login

[youngling@ada.iiit.ac.in](mailto:youngling@ada.iiit.ac.in)

or [youngling@10.4.24.24](mailto:youngling@10.4.24.24)

passwd: **BabyYodaBest**



# \$ partitions\_

```
[youngling@ada ~]$ sinfo
```

PARTITION	AVAIL	TIMELIMIT	NODES	STATE	NODELIST
short	up	6:00:00	2	idle	gnode[41-42]
long*	up	infinite	20	mix	gnode[01,11,23..]
long*	up	infinite	17	alloc	gnode[03-04,15,22..]
long*	up	infinite	21	drain	gnode[16,38..]
long*	up	infinite	2	maint	gnode[05-06]

# \$ nodestates\_

STATE	MEANING
IDLE	free to use
ALLOCATED	fully allocated
MIXED	have some CPUs free
MAINT	under maintenance
DRAIN	currently unavailable for use

these are the common ones. more are mentioned in the slurm pages

# \$ account-limits

## # Association Limits

```
sacctmgr show assoc \
  User="" Account=youngling \
  format=Account,GrpTRES,GrpTRESMins -p \
  | column -t -s "|";
```

## # Usage

```
scontrol show assoc_mgr \
  Account=youngling flags=assoc \
  | grep --color=auto -m 1 'GrpTRESMins'
```

# \$ storage

A youngling gets:

- 25GB on /home/youngling
- **50** (default) | 100 | 150 | 200 GB on /share3
- more storage can be allocated for special cases

use /share3 for long-term file storage like checkpoints, logs, etc.

Common datasets are curated and mounted at /share1/dataset



# \$ storage

In gnodeXY, you get write access on

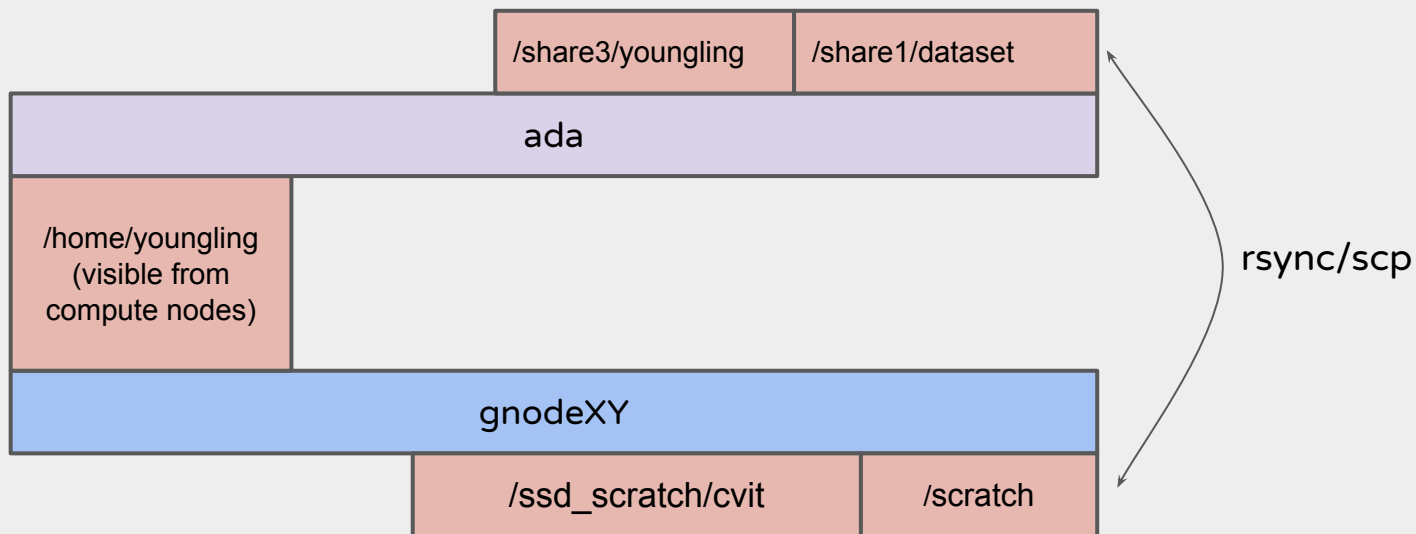
- `/ssd_scratch/cvit` - lasts 10 days, SSD, hence faster
- `/scratch` - lasts 10 days

for intermediate data storage when running models

# \$ storage-summary

	SPACE	QUOTA	VISIBILITY	PURPOSE	DELETION POLICY	BACKUP
headnode	/home	25 GB per user	head-node & gnode	small files, anaconda	-	Yes (daily)
	/share1	6TB (whole cvit)	only head-node	public datasets	-	No
	/share3	50 GB per user	only head-node	long-term storage	-	No
gnode	/scratch	2 TB (total)	only gnodeXY	temp hdd storage	10 days	No
	/scratch_ssd	960GB (total)	only gnodeXY	temp ssd storage for fast I/O	10 days	No

# \$ typical data workflow



# \$ warning

- Do not run heavy jobs on ada (head-node). Use it only to ssh in and gain access.
- Compile / Install packages only after gaining allocation on gnode. The OS mismatch will lead to your packages not working on gnodes (Ubuntu) if you install them on head-node (CentOS).
- /home/\$USER has a limit on number of files - 300k total
- Create folders nested under your name to identify what is yours. There are many concurrent users.
- 3 failed login attempts and you're locked out

[hands-on]

# \$ submit-job

- On ada, you submit “jobs” to gain access to a gnode
- Either use an interactive session (srun) or a headless session (sbatch)
- In either modes, mention the resource requirements followed by the command(s) to run on gaining access
- Avoid running interactive sessions on long partitions. Your idle time coding could be somebody else’s gpu time. Gain interactive shells on short partitions (6:00:00 limit)
- Your session will be terminated if you try to gain access without a GPU

# \$ module-load

- Environment modules allows users to set shell env variables needed to run programs

```
(base) [youngling@ada ~]$ module avail  
  
----- /opt/Modules/versions -----  
3.2.10  
  
----- /opt/Modules/3.2.10/modulefiles -----  
amber/18  
capnproto/0.6.1  
capnproto/0.7.0  
ceres-solver/1.14.0-165-gd7f428e  
cmake/3.15.2  
colmap/3.6-dev.2-15-g6b6e825  
cuda/10.0  
cuda/10.1  
cuda/10.2  
cuda/9.0  
cuda/9.1  
cudnn/7.1-cuda-9.1  
cudnn/7.3-cuda-10.0  
cudnn/7.6.4-cuda-9.0  
cudnn/7.6.5-cuda-10.2  
cudnn/7.6-cuda-10.0  
cudnn/7-cuda-10.0  
cudnn/7-cuda-9.0  
dot  
eigen/3.3.7  
ffmpeg/4.0.1  
ffmpeg/4.2.1  
freesurfer/6.0.0  
gflags/2.2.1  
gflags/2.2.2  
glog/0.3.5  
glog/0.4.0  
gromacs/2016.3  
gromacs/2016.3-plumed  
gromacs/2019  
gromacs/2019.3-plumed  
lammps/7Aug19  
lammps/7Aug19-v2  
leptonica/1.78.0  
matlab/R2019b  
matlab/R2020a  
mkl/2019.3.199  
mkl/2019.4.243  
module-git  
module-info  
modules  
mpich/3.3.1  
mrtrix/3.0  
namd/2.12  
namd/2.13  
null  
openblas/0.3.6  
opencv/4.1.2  
opencv/4.2.0  
openmesh/8.0  
openmpi/2.1.1  
openmpi/3.1.0  
openmpi/4.0.0  
openmpi/4.0.1-cuda10  
pcl/1.8.1  
plumed/2.5.2  
python/3.6.8  
python/3.7.4  
R/3.6.2  
singularity/2.5.2  
tbb/2018ui-debug  
tbb/2018ui-release  
tesseract/4.1.0  
use.own  
VTK/8.2.0
```

`module avail` - list all modules available

`module load <module-name>` - load module

> running-jobs\_

- srun
- sbatch



\$ srun ...

> srun

--pty

--account youngling

--time 4-00:00:00

/bin/bash -l

# interactive

# [<user>|research|sub]

# [days]-[hh]:[mm]:[ss]

# \$ simple-jobs\_

But defaults aren't usually nice to you! :(

```
> srun
  --pty                      # interactive
  --account youngling
  --partition long           # [short|long]
  --time 4-00:00:00          # [days]-[hh]:[mm]:[ss]
  --nodes 1                  # cpu cores
  --gres=gpu:0               # gpus
  /bin/bash -l
```

# \$ limits/upgrade

- cvit-queue is meant to be high availability and users expected to optimally use resources.
- Pay for your GPU usage by GPUMins. Your GPUMins is replenished **start of every month**. Use it wisely.
- GPUMins consumed = no. of GPUs used x total usage minutes (begins from allocation of gnode)
- initially each user is restricted to 1 GPU for 600 GPUmins/week. your usage pattern will be considered while upgrading.
  - prefer batch jobs to interactives

## \$ some-tips\_

prefer long-option arguments to the short-hands.  
the latter can be confusing.

--nodes    -N

--ntasks   -n

use the following proportions:

**1 gpu : 10 cores : 2-3GB mem-per-cpu**

(4 gpu: 40 cores : 120GB mem-per cpu ~ 1 node)

## \$ some-tips\_

**Keep your environment clean. Use conda/virtualenv**

Make a habit of using environments for your projects rather than installing packages at default /home/\$USER/bin

Avoids conflicting packages and is easier to clone on another system

## \$ some-tips\_

### .tar your large files


Always tar your large files (datasets, etc) as this significantly increases copy speed by eliminating file-creation overheads.

Untar them at the nodes.

**NOT** tar.gz (you don't want to compress uncompress).



Time to tinker





Create a small file in `ada:/share3/youngling`

Gain a bash shell through `srun`, try copying in and out of.

- `/scratch`
- `/ssd_scratch/cvit`



# \$ inspect

```
# inspect variables set by SLURM  
env | grep "SLURM"
```

```
# check properties of a specific job  
scontrol show jobs <jobid>
```

```
# search  
scontrol show jobs -o | grep "<regex>"
```

```
queue -u $USER # only your jobs
```

# \$ why am i not getting an alloc?

There could be many reasons shown:

- **Priority**

There are jobs waiting to be scheduled ahead of you in the queue.  
Ada's scheduling is FCFS.

- **Resources**

There are not enough resources to allocate to you.

# \$ why am i not getting an alloc?

- **AssocGrpGres**

Your group (cvit, ccnsb, research etc) is already using the max possible GPUs allocated to it. At max 60 GPUs can be used by all users belonging to the cvit user-group. Use `sinfo.x` command to check this.

- **AssocGrpCpuLimit | AssocGrpMemLimit:**

Same as above, except for CPU or Memory.

- **AssocGrpGresMins**

You have exhausted the *GPUMins* for the month. If you want it increased, contact cvit-admins.

# \$ when will I get an allocation?

```
scontrol show job <JobID> | grep -o StartTime=[^\ ]*
```

# \$ sbatch

can run headless.

scripts:

- setup to save stdout/stderr -> experiments are logged.
- commands are saved and documented.

mail notifications when complete.

**sbatch is how you can and should run distributed jobs.**

# \$ sbatch-script-skeleton

```
#!/bin/bash
```

```
# SBATCH ...
```

```
module load ... #
```

```
rsync ... or scp .. # Copy-In Data
```

```
function _export {<save-checkpoint-work>}
```

```
trap "_export" SIGXXX
```

```
# Code running
```

```
wait
```

```
_export
```

# \$ sbatch

```
#!/bin/bash
#SBATCH --account youngling
#SBATCH --nodes 1
#SBATCH --cores 8
#SBATCH --gres=gpu:1
#SBATCH --mem-per-cpu=2G
#SBATCH --time=1-00:00:00
#SBATCH --mail-type=END
#SBATCH --mail-user=<mail-id>

# Below command echoes commands as they're
# executed.
set -x;

echo "Running on gnode" >> ~/<dir>/log.txt
```

# \$ fault-tolerance > saving | resuming

```
#SBATCH --signal=B:TERM@900
```

```
CHECKPOINTS=...
```

```
REMOTE_DIR=...
```

```
function _export {  
    ssh $USER@ada "mkdir -p $REMOTE_DIR/"          # create-remote-directory  
    rsync -rz                                       \ # rsync  
        $CHECKPOINTS/checkpoint_{best,last}.pt    \ # src  
        ada:$REMOTE_DIR/                          # dest  
}
```

```
## We trap the SIGHUP signal and register a handler  
## which saves computed checkpoint  
## in this case
```

```
trap "_export" TERM
```

```
scancel <JOB-ID> - cancel currently running jobs
```




\$ more interactive-stuff

# remote-port-forwarding

Many a times, during prototyping it's convenient to bind ports of a gnode to your local-machine.

e.g:

- running a jupyter-notebook on ada
- running tensorboard on ada



# \$ remote-port-forwarding > jupyter

```
jupyter-notebook      \  
  --ip localhost      \  
  --port 8888         \  
  & # background
```

ssh

```
-N # Do not execute a remote command  
-f # Background, so we can proceed  
-R # Remote Port Forwarding  
  8888 # your-machine[bind_address:port]  
  :localhost:8888 # gnode[bind_address:port]  
  username@10.2.16.xx # where to ssh and set-this-up.
```



tensorboard? sftp?

# local-port forwarding

gnodeXY



app/  
store

Often, it makes sense to do the other way around, say for logging setups like visdom

This way, even if your job exited, your logs are saved on your local machine for you to inspect and won't have to fight other people for the accessing same node.

# \$ local-port-forwarding

```
# default - localhost:8097
```

```
[youngling@gnodexx]$ tensorboard --logdir=expt_1.0 &
```

```
[youngling@gnodexx]$
```

```
ssh
```

```
-N # Do not execute a remote command
```

```
-f # Background, so we can proceed
```

```
-L # Local Port Forwarding
```

```
8097:localhost:8097
```

```
username@ip # ip given by vpn(tun0). Starts with 10.2...
```

# \$ tips & tricks

- Mosh : handles connection timeouts gracefully. Instead of ``ssh username@ip`` -> ``mosh username@ip``
- tmux : run multiple shell instances without losing state. Will help when training models over hours/days
- You can use sh/bash/zsh on ada. Call the executable accordingly.
- If you want to run matlab or other interactive apps on ada, ssh with -X flag and load corresponding modules

# \$ tmux - terminal multiplexer

- Basic commands

- `tmux new -s <session-name>` # Creates new tmux session
- `tmux a -t <session-name>` # Attach to a running session
- `prefix(ctrl+b) + d` # Detach from session
- `tmux ls` # List all sessions
- `tmux kill-ses -t <session-name>` # Kill session

For more commands - `man tmux`

You can also find tmux cheat sheets online - [Link](#)

Homework(optional): tmux on ada (headnode) is old. Provide a build script so we can put one in cvit-contrib/modules for everyone.



# \$ tips & tricks

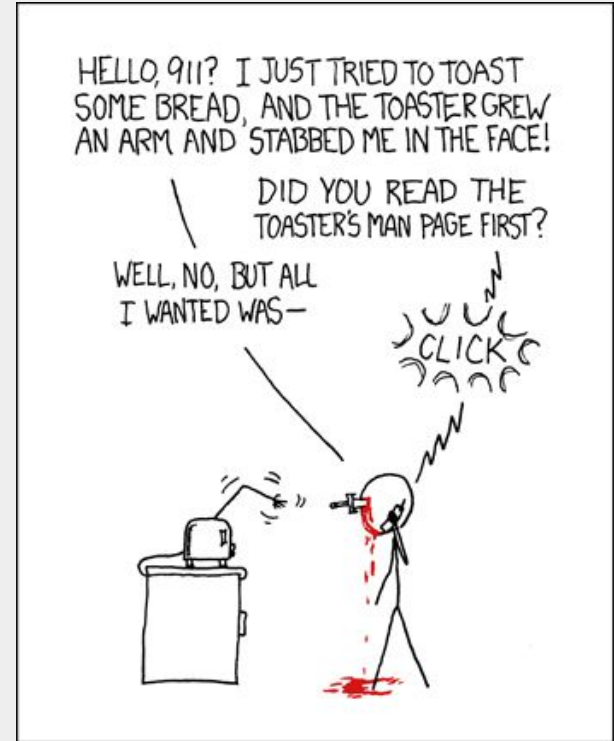
CUDA memory error? Not the end of the world.

Try :

- Make sure you are not performing unnecessary gradient computations
- Reduce batch size
- Acquire more GPUs and parallelize across them
  - > `model = nn.DataParallel(model)`
  - > `device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")`
  - > `model.to(device)`

# \$ man when in doubt

- all commands have extensive man pages. Get familiar with it.
- SLURM docs is your friend - [Slurm Workload Manager - Man Pages](#)
- you will run into errors, you have to send a paste of the log. [ix.io](#) is recommended. pipe error/output from ada shell and mail admins.



## \$ other-resources\_

- IIIT Ada User Guide - [http://hpc.iiit.ac.in/wiki/index.php/Ada\\_User\\_Guide](http://hpc.iiit.ac.in/wiki/index.php/Ada_User_Guide)
- Slurm docs - <https://slurm.schedmd.com/overview.html>
- Video tutorial by cvit-admins - [Video link](#)
- Sample sbatch script - [Github Gist](#)

For any issues, mail cvit-admins - [cvit-sudo@googlegroups.com](mailto:cvit-sudo@googlegroups.com)

# \$ Homework

- Figure out how to use tensorboard for your models (pytorch uses tensorboardx...already installed on youngling)
- Parallelize a CNN model from previous class using the data parallelism method you have studied. See if losses change with increase in batch-size (1, 2, 4, 8..)

Use youngling to experiment until you receive separate ada accounts.

Thank you!